

	<b>Scratch Install fitsstore</b> with mod_wsgi (Apache)
	<b>Kenneth Anderson</b> Paul Hirst
	Science User Support Department
	v1.1 – 30 May 2018

## Revision History

v1.0 – 1 March 2018      Kenneth Anderson  
v1.1 – 30 May 2018      Kenneth Anderson

**Document ID: GOA-PROC-101\_NewInstallFitsstoreApache**

## Document Purpose

Instructions for a scratch installation of fitsstore with mod\_wsgi on Apache web server.

## Intended Audience

Science User Support Department

## Table of Contents

1. Front Matter .....	1
2. Initialize Server and 'fitsdata' account.....	2
3. Setup the web server .....	3
4. Set up and Start Postgresql Database Server .....	5
5. Detailed Revision History .....	9

## 1. Front Matter

This is a prescription to install fitsstore under Apache running with mod\_wsgi. mod\_wsgi is the replacement for Apache mod\_python. You will need a user account on a machine to be designated as a fits server. For example, at Gemini-South, we do development on a development server, sbffits-dev-lv1.

Your user account will need `sudo -all` authorities. ITS will need to set up this account for you.

You will also need to be set up to use the **fitsdata** account on these servers. The **fitsdata** account will have very limited privileges -- this is the account Apache uses in the run environment.

I would recommend having two windows open, one logged in as **fitsdata** and one as your own account (w/ sudo). This way, it will be easy to switch from one to the other as the instructions will specify. (Some things will need sudo authorities, others will not).

These instruction pertain to a "clean install" on a CentOS7 machine; a number of packages need to be installed (system install) up front. A good amount of the following instruction set could potentially be placed into a shell script and executed in bulk. For example, all sudo pip install directives could be scripted.

## 2. Initialize Server and 'fitsdata' account

### Install and set up fitsdata account

Log in to *your* account (with sudo privileges) and begin on a command line:

```
$ sudo yum -y install httpd httpd-devel mod_ssl openssl
$ sudo yum -y install httpd postgresql postgresql-server postgresql-devel
```

### On AWS CENTOS7

Install EPEL repository (included by default on ITS config)

```
$ sudo yum install -y epel-release
$ sudo yum -y install gcc gcc-c++ python-pip python-devel gcc-gfortran cfitsio-devel
```

### Python package installation with pip

```
$ sudo pip install --upgrade pip
$ sudo pip install mod_wsgi==4.4.21
$ sudo pip install psycpg2
$ sudo pip install sqlalchemy
$ sudo pip install pyyaml
$ sudo pip install jinja2
$ sudo pip install pyfits
$ sudo pip install pywcs
$ sudo pip install dateutils
$ sudo pip install requests
$ sudo pip install matplotlib
$ sudo pip install scipy
$ sudo pip install pandas
$ sudo pip install astropy
```

If using s3 (for dev purposes, probably not)

```
$ sudo pip install boto3
```

### Set up the 'fitsdata' account

```
$ sudo /usr/sbin/groupadd -g 5179 fitsdata
$ sudo /usr/sbin/useradd -c 'FITS data' -u 5179 -g 5179 fitsdata
$ sudo passwd fitsdata
```

If this is a gemini fits server, edit /etc/group and add fitsdata

```
$ sudo vi(or emacs) /etc/group
geminidata:x:502:fitsdata
```

### Set up FitsStorage and DRAGONS

```
$ sudo mkdir /opt/FitsStorage
$ sudo chown fitsdata:fitsdata /opt/FitsStorage/
$ cd /opt/FitsStorage
$ svn checkout http://scisoft/svn/FitsStorage/trunk .
```

Note: If you forget to do this as 'fitsdata', and svn checkout with your sudo account (it happens!), you can chown the directory recursively.

```
$ sudo chown -R fitsdata:fitsdata FitsStorage/
```

### Installing DRAGONS

Get DRAGONS from GitHub.

Note: git pull will place 'DRAGONS' directly on cwd

```
$ cd /opt
$ sudo git clone https://github.com/GeminiDRSoftware/DRAGONS.git
```

Now, under /opt you should have directories with owners that look like,

```
drwxr-xr-x 11 fitsdata  4096 Feb  8 16:24 FitsStorage/
drwxr-xr-x 13 <youract> 4096 Feb  7 17:45 DRAGONS
```

### Install fitsverify

As part of file validation, FitsStorage makes use of the executable, "fitsverify".

This is an executable built from the cfitsio library. You can find this executable on current production Fits servers in /opt. Add this to /opt on the new server and make the owner 'fitsdata'. Your /opt directory should then look something like,

```
/opt:
drwxr-xr-x 12 <youract>  root 4096 May 11 16:22 DRAGONS
drwxr-xr-x 12 fitsdata  fitsdata 4096 May  8 15:58 FitsStorage
drwxr-xr-x  2 <youract>  root 4096 May 20 15:26 fitsverify
```

It likely will be true that the 'fitsverify' binary will need to be compiled from source if the OS is something other than CentOS7.

## **3. Setup the web server**

Your sudo account will be used to set up the web server.

The following is a shell "install script" that will provide necessary configuration for the Apache server using mod\_wsgi. This is installed with mod\_wsgi-express. You can copy and paste this into a file (suggestion: > /tmp/install\_script.sh) and run it.

(Modify the variables at the beginning of the following script if needed)

---

```
# NOTE!!! IF YOU MODIFY THE FOLLOWING, THE CHANGES NEED TO BE
# TRANSLATED TO /etc/systemd/system/fits-httpd.service!!

FITSSTORAGE_ROOT=/opt/FitsStorage
DRAGONS_ROOT=/opt/DRAGONS
MODWSGI_DIR=/opt/modwsgi-default

# Explanation of options that we are setting in MODWSGI_OPTS
# --server-root          where to store server setup information
# --access-log           Enable access log (off by default)
# --url-alias            Alias URL for static files (one per alias)
# --python-path         Additional content for PYTHONPATH

MODWSGI_OPTS="\
--server-root $MODWSGI_DIR \
--port 80 \
--user apache --group apache \
--access-log \
--url-alias /static $FITSSTORAGE_ROOT/htmldocroot \
--url-alias /favicon.ico $FITSSTORAGE_ROOT/htmldocroot/favicon.ico \
--url-alias /robots.txt $FITSSTORAGE_ROOT/htmldocroot/robots.txt \
--python-path $FITSSTORAGE_ROOT \
```

```
--python-path $DRAGONS_ROOT \
"
mkdir ${MODWSGI_DIR}

# The extra argument is the WSGI entry point

mod_wsgi-express setup-server\
$MODWSGI_OPTS $FITSSTORAGE_ROOT/fits_storage/wsgihandler.py

# And set let apache handle that dir from now on
chown -R apache:apache $MODWSGI_DIR

EOF
```

---

### **Run the install script to set up mod\_wsgi**

After saving the above install script (here we have written it to /tmp), execute the script.

```
$ sudo bash /tmp/install_script.sh
```

(If you have written the file somewhere else, execute that, of course.)

To run the web server from systemd

```
$ sudo cp /opt/FitsStorage/otherfiles/etc_systemd_system_fits-httpd.service
/etc/systemd/system/fits-httpd.service
```

```
$ sudo cp /opt/FitsStorage/otherfiles/etc-sysconfig-httpd /etc/sysconfig/httpd
```

Edit the config file generated by the install\_script.sh. This configuration file will be /opt/modwsgi-default/httpd.conf.

Add these lines by the end of the "Alias" list. Then edit to File \* for permissions:

```
AliasMatch '/(.*\.html)' '/opt/FitsStorage/htmldocroot/$1'
AliasMatch '/(.*\.css)' '/opt/FitsStorage/htmldocroot/$1'
AliasMatch '/(.*\.js)' '/opt/FitsStorage/htmldocroot/$1'
```

Which should have httpd.conf looking like:

```
[ ... ]
Alias '/favicon.ico' '/opt/FitsStorage/htmldocroot/favicon.ico'

<Directory '/opt/FitsStorage/htmldocroot'>
<Files '*'>
    Require all granted
</Files>
</Directory>

AliasMatch '/(.*\.html)' '/opt/FitsStorage/htmldocroot/$1'
AliasMatch '/(.*\.css)' '/opt/FitsStorage/htmldocroot/$1'
AliasMatch '/(.*\.js)' '/opt/FitsStorage/htmldocroot/$1'

<IfDefine MOD_WSGI_VERIFY_CLIENT>
<Location '/'>
    SSLVerifyClient require
    SSLVerifyDepth 1
```

```
</Location>
</IfDefine>
[ ... ]
```

In the httpd.conf file, find LimitRequestBody and comment it out:

```
#LimitRequestBody 200000000
```

Comment out all the \_default\_ VirtualHost definitions.

E.g.,

```
<IfDefine !MOD_WSGI_VIRTUAL_HOST>
<IfVersion < 2.4>
NameVirtualHost *:80
</IfVersion>
#<VirtualHost _default_:80>
#</VirtualHost>
</IfDefine>
```

Be sure to comment a full element block – do not create badly formed xml. The next virtual host element should look like,

```
<IfDefine MOD_WSGI_VIRTUAL_HOST>

<IfVersion < 2.4>
NameVirtualHost *:80
</IfVersion>
#<VirtualHost _default_:80>
#<Location />
#Order deny,allow
#Deny from all
#<IfDefine MOD_WSGI_ALLOW_LOCALHOST>
#Allow from localhost
#</IfDefine>
#</Location>
#</VirtualHost>
```

### **Edit /opt/modwsgi-default/handler.wsgi**

Set "entry point". The default install of modwsgi-default sets the entry point in the handler.wsgi file to be

```
entry_point = '/opt/modwsgi-default/default.wsgi'
```

Set this to use the FitsStorage wsgi handler.py:

```
entry_point = '/opt/FitsStorage/fits_storage/wsgihandler.py'
```

## 4. Set up and Start Postgresql Database Server

### Set up log and backup directories

```
$ sudo mkdir /data/logs
$ sudo chown fitsdata /data/logs
$ sudo mkdir /data/backups
$ sudo chown fitsdata /data/backups
$ sudo mkdir /data/upload_staging
$ sudo chown fitsdata /data/upload_staging
$ sudo chmod o+rw /data/upload_staging
$ sudo mkdir /data/z_staging
$ sudo chown fitsdata /data/z_staging
```

If using s3:

```
$ sudo mkdir /data/s3_staging
$ sudo chown fitsdata /data/s3_staging
```

Apache needs access to this to download preview images now

```
$ sudo chmod o+rw /data/s3_staging
```

### Set up the postgresql data directory

```
$ sudo mkdir /data/pgsql_data
$ sudo chown postgres:postgres /data/pgsql_data
$ sudo cp /lib/systemd/system/postgresql.service /etc/systemd/system/
```

Edit the postgresql.service file just placed in /etc/systemd/system/ and add the PGDATA environment var:

```
$ sudo vi (emacs) /etc/systemd/system/postgresql.service
```

```
Environment=PGDATA=/data/pgsql_data
```

(save & exit)

### Initiate the database

```
$ sudo systemctl daemon-reload
$ sudo postgresql-setup initdb
$ sudo systemctl start postgresql.service
$ sudo su - postgres
$ /usr/bin/createuser --no-superuser --no-creatorole --createdb fitsdata
$ /usr/bin/createuser --no-superuser --no-creatorole --no-createdb apache
```

Edit the postgresql configuration and set the effective cache size

```
$ emacs (vi) /data/pgsql_data/postgresql.conf
```

```
#shared_buffers = 256MB
# 25% of memory if memory > 1GB
```

```
effective_cache_size = 8000MB
```

Save and exit.

If there will be remote access to this database (eg for eavedropping), edit postgresql.conf and set,

```
listen_addresses = '*'
```

And edit /data/pgsql\_data/pg\_hba.conf . At end of file, add a comment line to say what you're adding. E.g.,

```
host fitsdata all 10.1.5.47/32 trust exit
```

Save the file and exit.

### **Start Postgresql**

```
$ sudo systemctl restart postgresql.service
```

Need to set postgresql and httpd to start at boot

```
$ sudo systemctl enable postgresql.service
$ sudo systemctl enable fits-httpd.service
```

Now get a shell logged in as 'fitsdata'

```
$ cd /opt/FitsStorage/fits_storage
$ emacs (vi) fits_storage_config.py
# Set values
# This is a setup for a fitsstore server, not the archive. Set that boolean:
```

```
use_as_archive = False
```

```
# Set the storage root to point to a data source. For testing, place test
# data under /data/
```

```
storage_root = '/data/gemini_data'
```

```
# On production servers, this will point to dataflow:
```

```
storage_root = '/sci/dataflow'
```

(save and exit)

### **Create db and set up tables.**

```
$ /usr/bin/createdb fitsdata
$ python scripts/create_tables.py
```

### **Should now be good to go.**

- Check with ITS if you need a whitelisted local mail relay on it
- Check with ITS that the development/operations status of the machine

To run jobs from system

```
$ sudo cp /opt/FitsStorage/otherfiles/etc_systemd_system_fits-service_ingest_queue1.service  
/etc/systemd/system/fits-service_ingest_queue1.service  
$ sudo systemctl start fits-service_ingest_queue1  
$ sudo systemctl enable fits-service_ingest_queue1
```

Copy again and edit for siq2 if you want two of them

- Same for 2 if you want it
- Repeat for export queue and api



## 5. Detailed Revision History

v1.0	1 March 2018	Kenneth Anderson
v1.1	30 May 2018	Kenneth Anderson